

CS501- Week 10 Homework 1: Machine Learning on Kubernetes

Name: Feven Belay Araya

ID: 20027

https://hc.labnet.sfbu.edu/~henry/sfbu/course/cloud_computing/genai/slide/exercise_kubernetes.html

Q2 ==> Machine Learning on Kubernetes

Machine Learning on Kubernetes

Creating and uploading necessary files in GCP- Cloud Shell Terminal

1. Start minikube in Google Cloud Platform

```
fba8584@cloudshell:~ (sfbu-cs571-414319)$ minikube start
* minikube v1.32.0 on Debian 11.9 (amd64)
- MINIKUBE_FORCE_SYSTEMD=true
- MINIKUBE_HOME=/google/minikube
- MINIKUBE_WANTUPDATENOTIFICATION=false
* Using the docker driver based on existing profile
* Starting control plane node minikube in cluster minikube
* Pulling base image ...
* Updating the running docker "minikube" container ...

X Docker is nearly out of disk space, which may cause deployments to fail! (95% of capacity). You can pass '--force' to skip this check.
* Suggestion:

  Try one or more of the following to free up space on the device:

  1. Run "docker system prune" to remove unused Docker data (optionally with "-a")
  2. Increase the storage allocated to Docker for Desktop by clicking on:
     Docker icon > Preferences > Resources > Disk Image Size
  3. Run "minikube ssh -- docker system prune" if using the Docker container runtime
* Related issue: https://github.com/kubernetes/minikube/issues/9024

* Preparing Kubernetes v1.28.3 on Docker 24.0.7 ...
- kubelet.cgroups-per-qos=false
- kubelet.enforce-node-allocatable=""
* Verifying Kubernetes components...
- Using image gcr.io/k8s-minikube/storage-provisioner:v5
* Enabled addons: storage-provisioner, default-storageclass
* Done! kubect1 is now configured to use "minikube" cluster and "default" namespace by default
```

2. Create requirements.txt file using the following command

- nano requirements.txt

```
fba8584@cloudshell:~ (sfbu-cs571-414319)$ nano requirements.txt
```

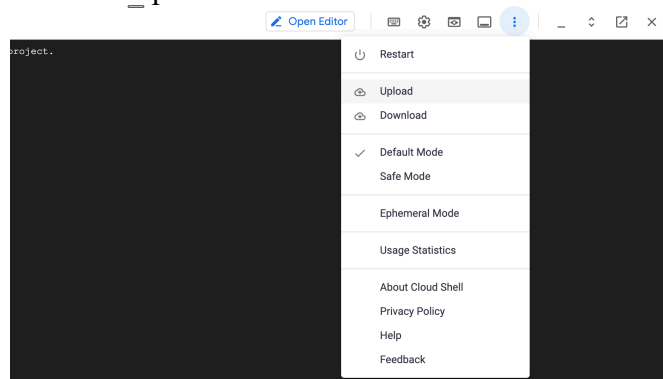
Then enter the following contents

```
Flask==1.1.1
unicorn==19.9.0
itsdangerous==1.1.0
Jinja2==2.10.1
MarkupSafe==1.1.1
Werkzeug==0.15.5
numpy==1.19.5 # Adjusted to a version before np.float deprecation
scipy>=0.15.1
scikit-learn==0.24.2 # Ensure compatibility with numpy version
matplotlib>=1.4.3
pandas>=0.19
flasgger==0.9.4
```

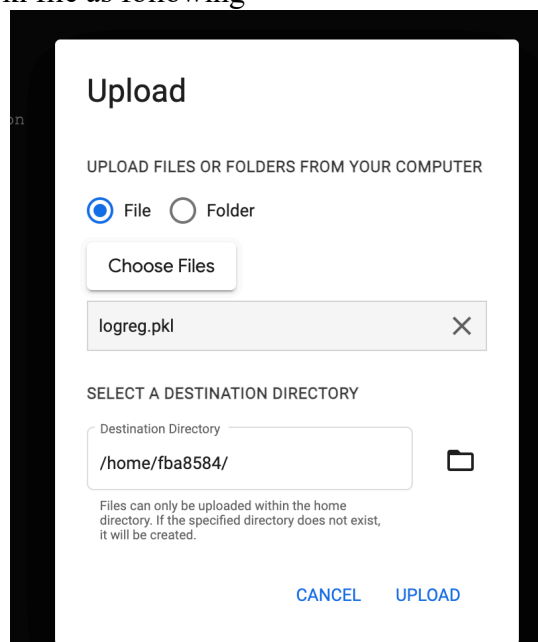


```
GNU nano 5.4
Flask==1.1.1
gunicorn==19.9.0
itsdangerous==1.1.0
Jinja2==2.10.1
MarkupSafe==1.1.1
Werkzeug==0.15.5
numpy==1.19.5 # Adjusted to a version before np.float deprecation
scipy>0.15.1
scikit-learn==0.24.2 # Ensure compatibility with numpy version
matplotlib>1.4.3
pandas>0.19
flasgger==0.9.4
```

3. Upload logreg.pkl file by clicking the three dots in the top-right part of the Cloud Shell Terminal and then choose upload



Then upload the logreg.pkl file as following



4. Create flask_api.py file using the command
- *nano flask_api.py*

```
fba8584@cloudshell:~ (sfbu-cs571-414319) $ nano flask_api.py
```

Then enter the following contents

```
# -*- coding: utf-8 -*-
"""
Created on Mon May 25 12:50:04 2020

@author: pramod.singh
"""

from flask import Flask, request
import numpy as np
import pickle
import pandas as pd
from flasgger import Swagger

app = Flask(__name__)
Swagger(app)

pickle_in = open("logreg.pkl", "rb")
model = pickle.load(pickle_in)

@app.route('/')
def home():
    return "Welcome to the Flask API!"

@app.route('/predict', methods=["GET"])
def predict_class():
    """Predict if Customer would buy the product or not.
    ---
    parameters:
      - name: age
        in: query
        type: number
        required: true
      - name: new_user
        in: query
        type: number
        required: true
      - name: total_pages_visited
        in: query
        type: number
        required: true
    responses:
      200:
        description: Prediction
    """
    age = int(request.args.get("age"))
    new_user = int(request.args.get("new_user"))
    total_pages_visited = int(request.args.get("total_pages_visited"))
    prediction = model.predict([[age, new_user, total_pages_visited]])
    return "Model prediction is " + str(prediction)

@app.route('/predict_file', methods=["POST"])
def prediction_test_file():
    """Prediction on multiple input test file.
```

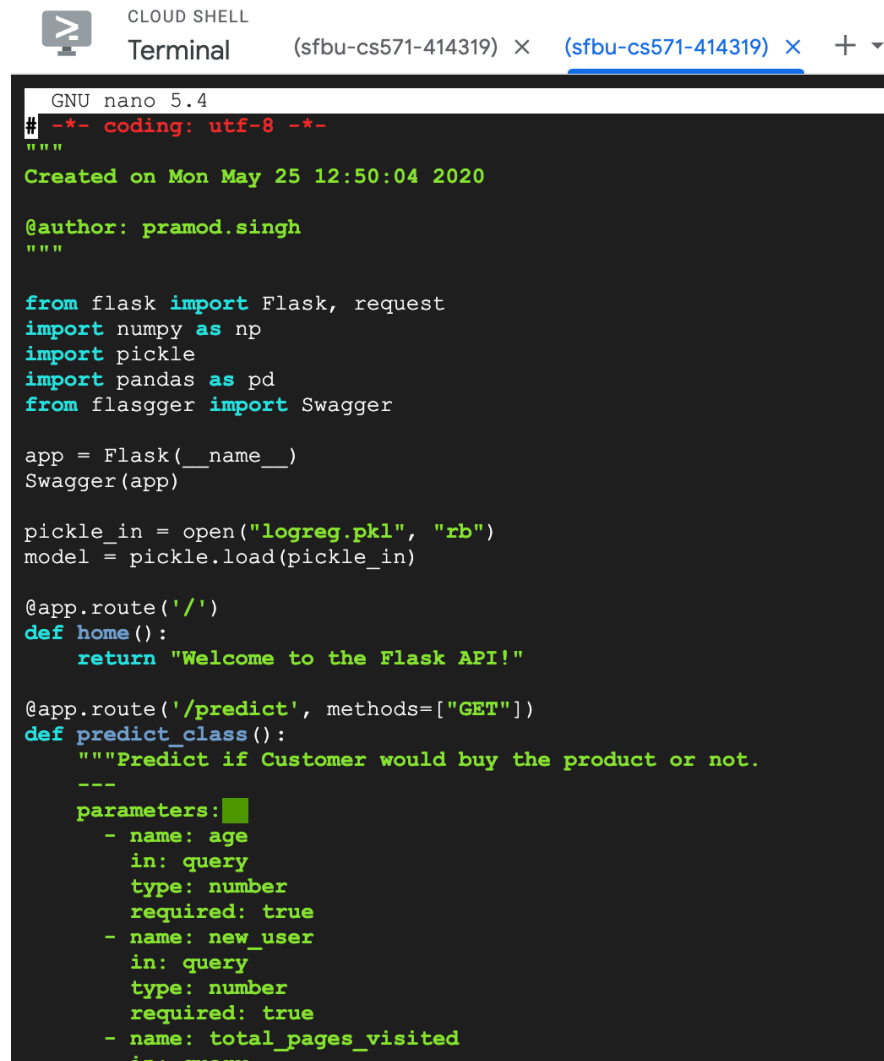
```

---
parameters:
  - name: file
    in: formData
    type: file
    required: true
responses:
  200:
    description: Test file Prediction
"""

df_test = pd.read_csv(request.files.get("file"))
prediction = model.predict(df_test)
return str(list(prediction))

if __name__ == '__main__':
    app.run(debug=True, host='0.0.0.0', port=5000)

```



The screenshot shows a Cloud Shell terminal window with a nano editor open. The code in the editor is a Python script for a Flask API. It includes imports for Flask, request, numpy, pickle, pandas, and Swagger. The script defines a Flask app, loads a pickle model, and sets up two routes: a home page and a predict endpoint. The predict endpoint uses a Swagger decorator and a predict_class function that returns a list of predictions. The code is partially visible, showing the beginning of the predict_class function's parameters section.

```

GNU nano 5.4
# -*- coding: utf-8 -*-
"""
Created on Mon May 25 12:50:04 2020

@author: pramod.singh
"""

from flask import Flask, request
import numpy as np
import pickle
import pandas as pd
from flasgger import Swagger

app = Flask(__name__)
Swagger(app)

pickle_in = open("logreg.pkl", "rb")
model = pickle.load(pickle_in)

@app.route('/')
def home():
    return "Welcome to the Flask API!"

@app.route('/predict', methods=["GET"])
def predict_class():
    """Predict if Customer would buy the product or not.
    ---
    parameters:
      - name: age
        in: query
        type: number
        required: true
      - name: new_user
        in: query
        type: number
        required: true
      - name: total_pages_visited
        in: query

```

```

GNU nano 5.4
def predict_class():
    """Predict if Customer would buy the product or not.
    ---
    parameters:
      - name: age
        in: query
        type: number
        required: true
      - name: new_user
        in: query
        type: number
        required: true
      - name: total_pages_visited
        in: query
        type: number
        required: true
    responses:
      200:
        description: Prediction
    """
    age = int(request.args.get("age"))
    new_user = int(request.args.get("new_user"))
    total_pages_visited = int(request.args.get("total_pages_visited"))
    prediction = model.predict([age, new_user, total_pages_visited])
    return "Model prediction is " + str(prediction)

@app.route('/predict_file', methods=["POST"])
def prediction_test_file():
    """Prediction on multiple input test file.
    ---
    parameters:
      - name: file
        in: formData
        type: file
        required: true
    responses:
      200:
        description: Test file Prediction
    """

```

```

@app.route('/predict_file', methods=["POST"])
def prediction_test_file():
    """Prediction on multiple input test file.
    ---
    parameters:
      - name: file
        in: formData
        type: file
        required: true
    responses:
      200:
        description: Test file Prediction
    """
    df_test = pd.read_csv(request.files.get("file"))
    prediction = model.predict(df_test)
    return str(list(prediction))

@app.route('/apidocs')
def api_docs():
    return "API Documentation goes here."

if __name__ == '__main__':
    app.run(debug=True, host='0.0.0.0', port=5000)

```

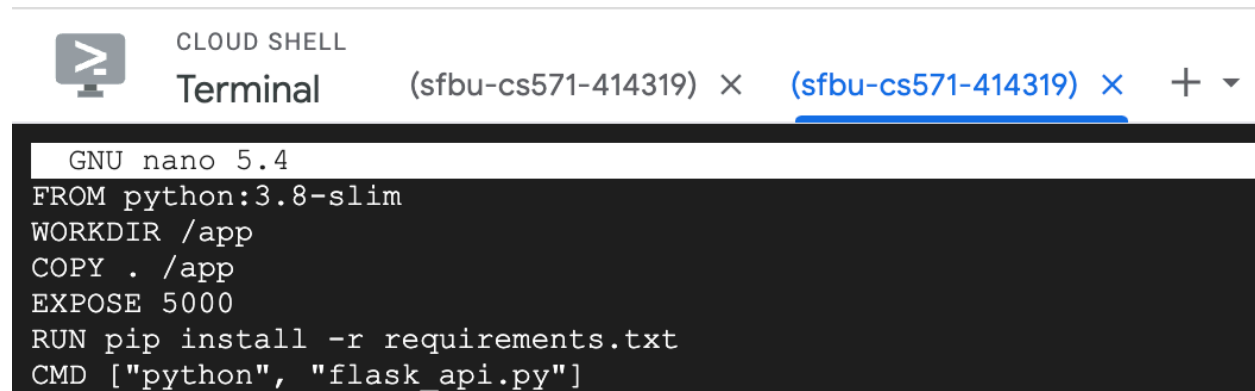
Step 4: Dockerfile

1. Create Dockerfile using command
 - `nano Dockerfile`

```
fba8584@cloudshell:~ (sfbu-cs571-414319) $ nano Dockerfile
```

Then enter the following content

```
FROM python:3.8-slim
WORKDIR /app
COPY . /app
EXPOSE 5000
RUN pip install -r requirements.txt
CMD ["python", "flask_api.py"]
```



The screenshot shows a Cloud Shell terminal window with the title 'Terminal (sfbu-cs571-414319)'. Inside the terminal, the GNU nano 5.4 editor is open, displaying the following content in a Dockerfile:

```
GNU nano 5.4
FROM python:3.8-slim
WORKDIR /app
COPY . /app
EXPOSE 5000
RUN pip install -r requirements.txt
CMD ["python", "flask_api.py"]
```

1. 'FROM python:3.8-slim'

- This line sets the base image for the Docker image you are creating. It tells Docker to start with the 'python:3.8-slim' image, which is an official Python image with Python 3.8 installed on it. The 'slim' version is a smaller version of the image that has fewer packages pre-installed, making the image size smaller.

2. 'WORKDIR /app'

- This instruction sets the working directory within the Docker container to `*/app*`. All subsequent commands will be executed in this directory within the container.

3. 'COPY . /app'

- This line copies everything from the current directory (on the host machine where you're running the Docker build command, indicated by the first `..`) into the `*/app` directory inside the Docker image (the second `*/app*`).

4. 'EXPOSE 5000'

- The 'EXPOSE' instruction informs Docker that the container listens on the specified network port at runtime. In this case, it tells Docker that the container will listen on port 5000. It's worth noting that this does not actually publish the port—it serves as documentation and is used by the 'docker run -p' command to map the container port to a port on the Docker host.

5. 'RUN pip install -r requirements.txt'

- This command tells Docker to run `pip install` inside the container, which will install the Python dependencies listed in the `requirements.txt` file. These dependencies are necessary for the Flask application to run correctly.

6. 'CMD ["python", "flask_api.py"]'

- This is the command that will be executed by default when the Docker container starts. In this case, it's telling Docker to run `flask_api.py` using Python. This is the Flask application you want to run inside the container.

Step 5: Running the Docker Container

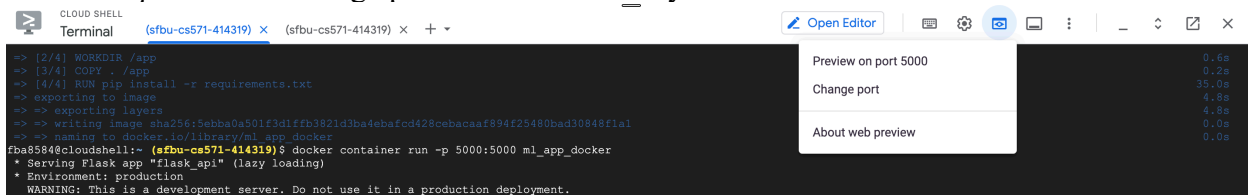
1. To build the docker image use the command
- `sudo docker build -t ml_app_docker .`

```
fba8584@cloudshell:~ (sfbu-cs-571-414319)$ docker build -t ml_app_docker .
[+] Building 45.2s (9/9) FINISHED
=> [internal] load build definition from Dockerfile
=> => transferring dockerfile: 150B
=> [internal] load metadata for docker.io/library/python:3.8-slim
=> => transferring context: 2B
=> [1/4] FROM docker.io/library/python:3.8-slim@sha256:72ae14e80c21f274f3111deb505d8fa64536fdf41b57f03930b3baf84d8b8d
=> => resolve docker.io/library/python:3.8-slim@sha256:72ae14e80c21f274f3111deb505d8fa64536fdf41b57f03930b3baf84d8b8d
=> => sha256:72ae14e80c21f274f3111deb505d8fa64536fdf41b57f03930b3baf84d8b8d 1.86kB / 1.86kB
=> => sha256:0a0d39c2b848b1b348f80ad80a43c4139469e6a49749153d2131e2a2a73f 1.37kB / 1.37kB
=> => sha256:04977f08fab13b02b303d547d2cc03e74e052ae9f9b0be6ae2ab90758fdae 6.97kB / 6.97kB
=> => sha256:8a1e25ce7c4f75e372e9884f8f7b1bedcfe4a7a7d452ebdb0alc7477c9a90345 29.12MB / 29.12MB
=> => sha256:1103112ebfc6e01c0f35f3586e5a39c6a9ffa32c1a362d4d5f20e3783c6fdd7 3.51MB / 3.51MB
=> => sha256:93d3f6d14ae5338f6f6394ed5946980d38c016a537a330f20921c5c7e3995a9 11.67MB / 11.67MB
=> => sha256:18dacb39e6d34eadfba0da78f1b3a5f3addfcd45ee854f6af9877b9ed5c4b1f 3.13MB / 3.13MB
=> => sha256:46996c1c5e3592977cd1c8454cf833bf486a5be36f71847794d97bac47a35f0 246B / 246B
=> => extracting sha256:8a1e25ce7c4f75e372e9884f8f7b1bedcfe4a7a7d452ebdb0alc7477c9a90345 1.7s
=> => extracting sha256:1103112ebfc6e01c0f35f3586e5a39c6a9ffa32c1a362d4d5f20e3783c6fdd7 0.2s
=> => extracting sha256:93d3f6d14ae5338f6f6394ed5946980d38c016a537a330f20921c5c7e3995a9 0.6s
=> => extracting sha256:46996c1c5e3592977cd1c8454cf833bf486a5be36f71847794d97bac47a35f0 0.3s
=> [internal] load build context
=> => transferring context: 29.86MB
=> [2/4] WORKDIR /app
=> [3/4] COPY . /app
=> [4/4] RUN pip install -r requirements.txt
=> exporting to image
=> exporting layers
=> writing image sha256:5ebba0a501f3d1ffb3821d3ba4ebafcd428cebacaa894f25480bad30848f1a1
=> naming to docker.io/library/ml_app_docker
0.0s
```

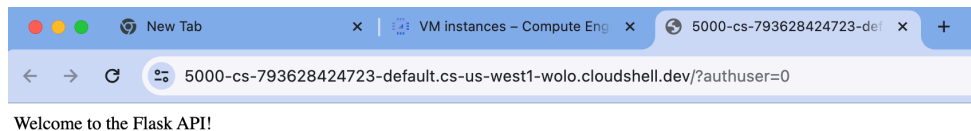
2. This command runs a Docker container from the ml_app_docker image:
- `docker container run -p 5000:5000 ml_app_docker`

```
fba8584@cloudshell:~ (sfbu-cs-571-414319)$ docker container run -p 5000:5000 ml_app_docker
* Serving Flask app "flask_api" (lazy loading)
* Environment: production
WARNING: This is a development server. Do not use it in a production deployment.
Use a production WSGI server instead.
* Debug mode: on
/usr/local/lib/python3.8/site-packages/sklearn/base.py:310: UserWarning: Trying to unpickle estimator LogisticRegression from version 0.23.2 when using version 0.24.2. This might le
ad to breaking code or invalid results. Use at your own risk.
warnings.warn(
* Running on http://0.0.0.0:5000/ (Press CTRL+C to quit)
* Restarting with stat
/usr/local/lib/python3.8/site-packages/sklearn/base.py:310: UserWarning: Trying to unpickle estimator LogisticRegression from version 0.23.2 when using version 0.24.2. This might le
ad to breaking code or invalid results. Use at your own risk.
warnings.warn(
* Debugger is active!
* Debugger PIN: 297-165-261
```

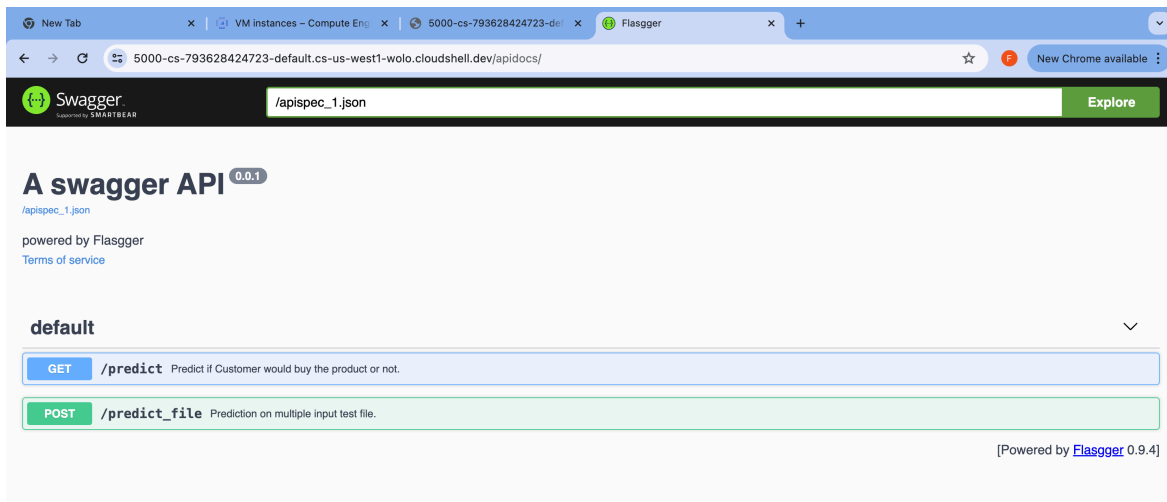
3. In the right-upper side of the terminal click the eye shaped button and then click *Preview on port 5000*. Change port if it is not 5000 by default.



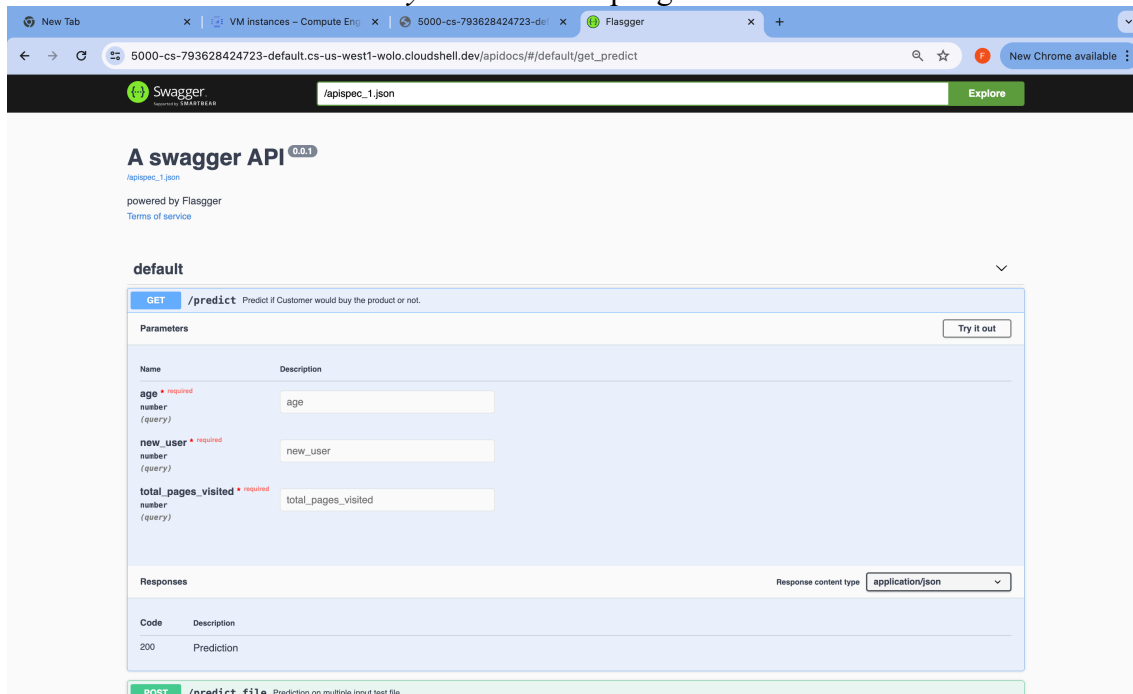
4. You will see this using the web preview.



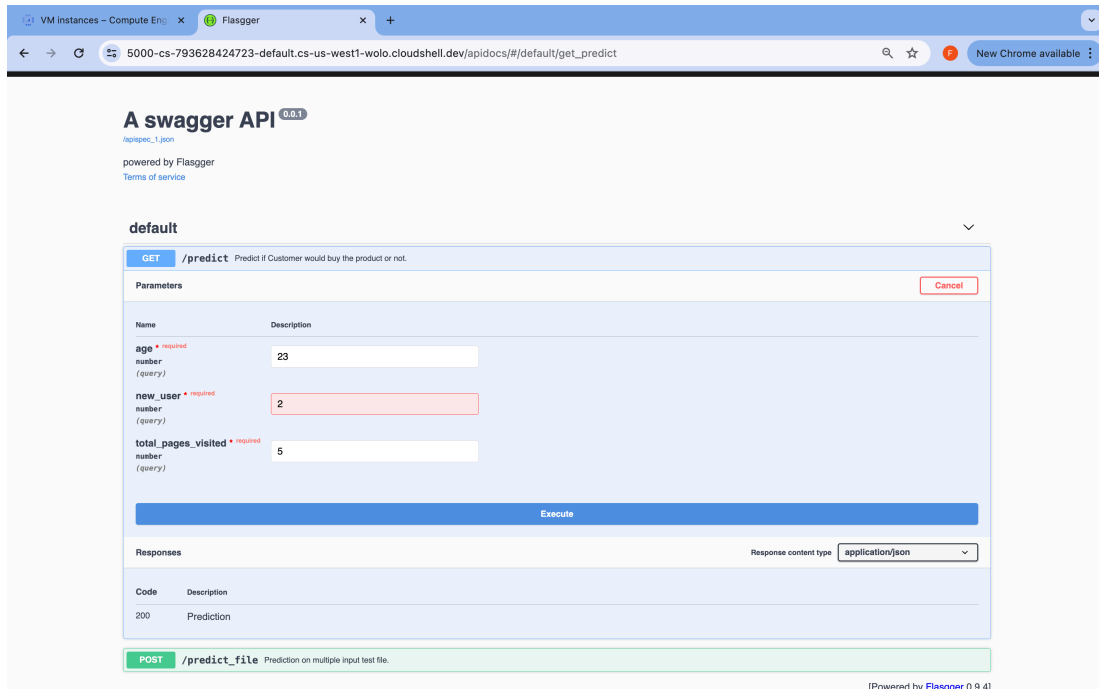
5. Add /apidocs/ at the end of the link to access the running ml- app as following
- There are two tabs GET and POST.



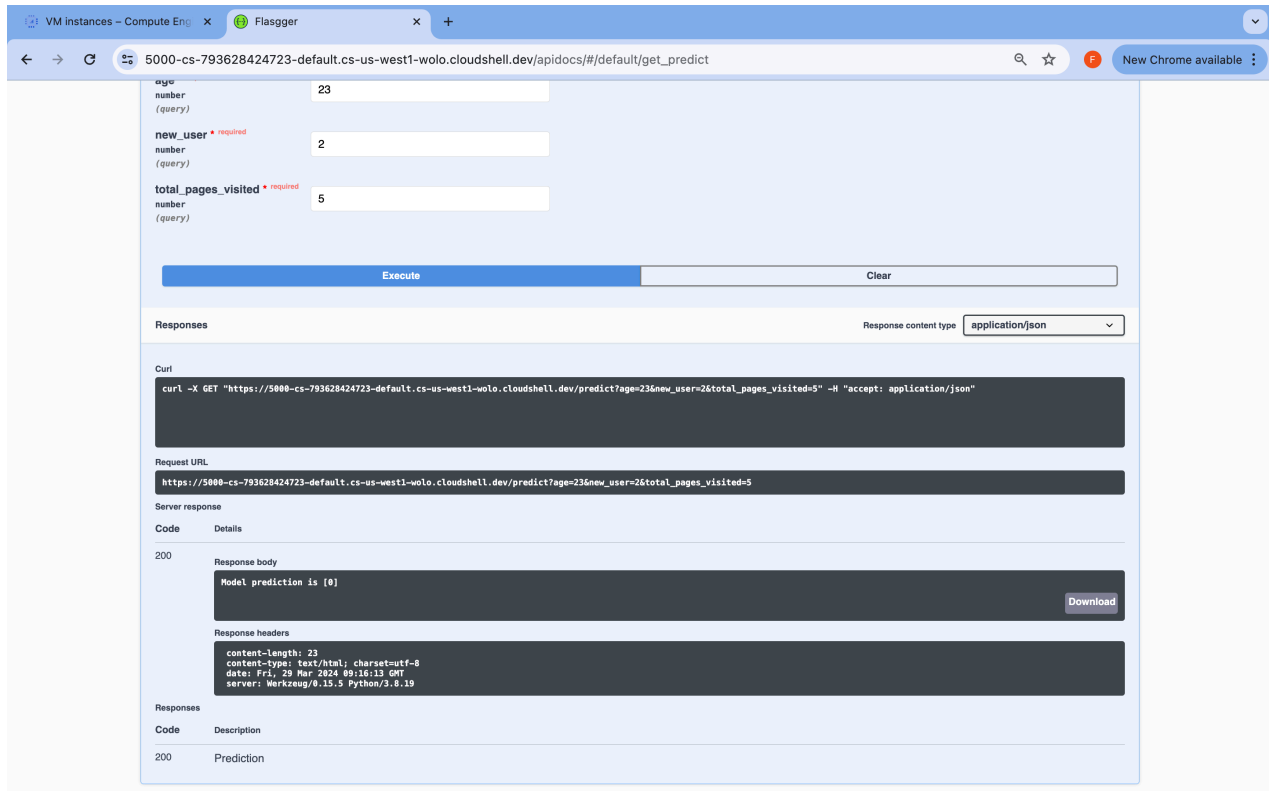
6. Click *GET* and then click *Try it out* in the top-right corner of the GET box.



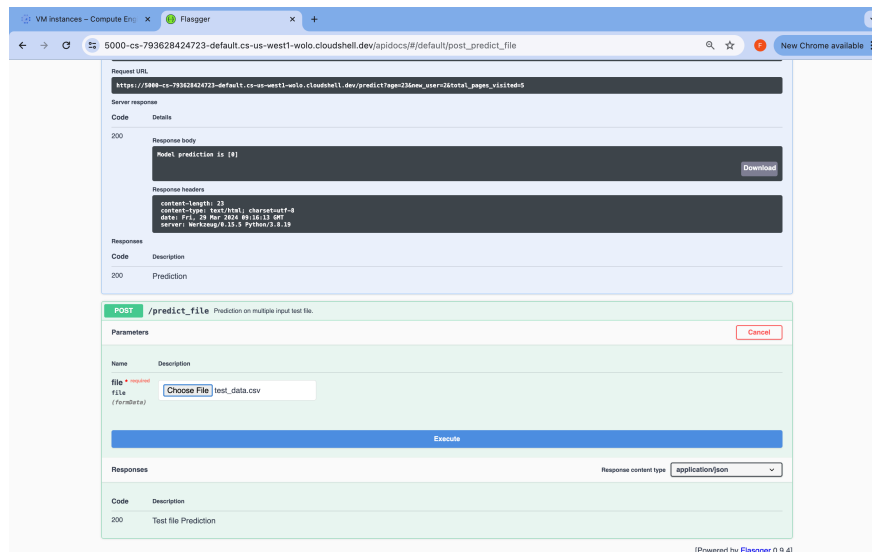
7. Fill values for the input parameters and then click Execute.



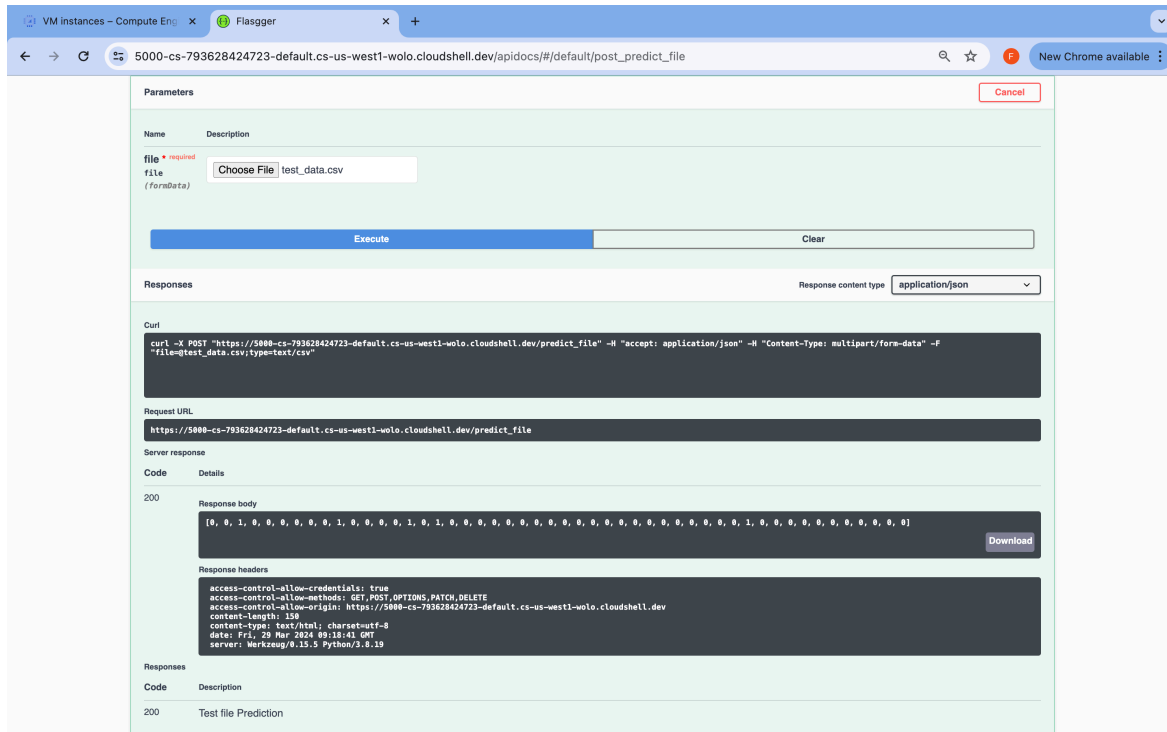
8. Upon the execution call, the request goes to the app, and predictions are made by the model.
 - The result of the model prediction is displayed in the Prediction section of the page as following



- The next prediction that can be done is for a group of customers (test data) via a post request.



10. Upload the test data file containing the same parameters in a similar order. The model would make the prediction, and the results would be displayed upon execute as following.



Step 6: Stopping/killing the running container

1. Use `docker ps` to list running Docker containers

```
fba8584@cloudshell:~ (sfbu-cs571-414319) $ docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
3671198b7a37	ml_app_docker	"python flask api.py"	About an hour ago	Up About an hour	0.0.0.0:5000->5000/tcp	ecstatic joliot

- The CONTAINER ID is given as 367119b87a37

2. Use the command

- `docker kill <CONTAINER ID>` to kill the running container as follows.

```
fba8584@cloudshell:~ (sfbu-cs571-414319) $ docker kill 367119b87a37
367119b87a37
fba8584@cloudshell:~ (sfbu-cs571-414319) $
```